

# Search Problems


## Theory and Applications



Prof. Dr. Eduardo Noronha  
Instituto Federal de Goiás (IFG)


# Today

- Agents that Plan Ahead
- Search Problems
- Uninformed Search Methods
  - Depth-First Search
  - Breadth-First Search
  - Uniform-Cost Search
- Implementation
- Applications



# Reflex Agents

- Reflex agents:
  - Choose action based on current percept (and maybe memory)
  - May have memory or a model of the world's current state
  - Do not consider the future consequences of their actions
  - Consider how the world IS
- Can a reflex agent be rational?




[Demo: reflex optimal (L2D1)]

[Demo: reflex optimal (L2D2)]


# Reflex Agent - Video

---




# Reflex Odd - Video

---




# Planning Agents

- Planning agents:
  - Ask “what if”
  - Decisions based on (hypothesized) consequences of actions
  - Must have a model of how the world evolves in response to actions
  - Must formulate a goal (test)
  - Consider how the world **WOULD BE**
- Optimal vs. complete planning
- Planning vs. replanning

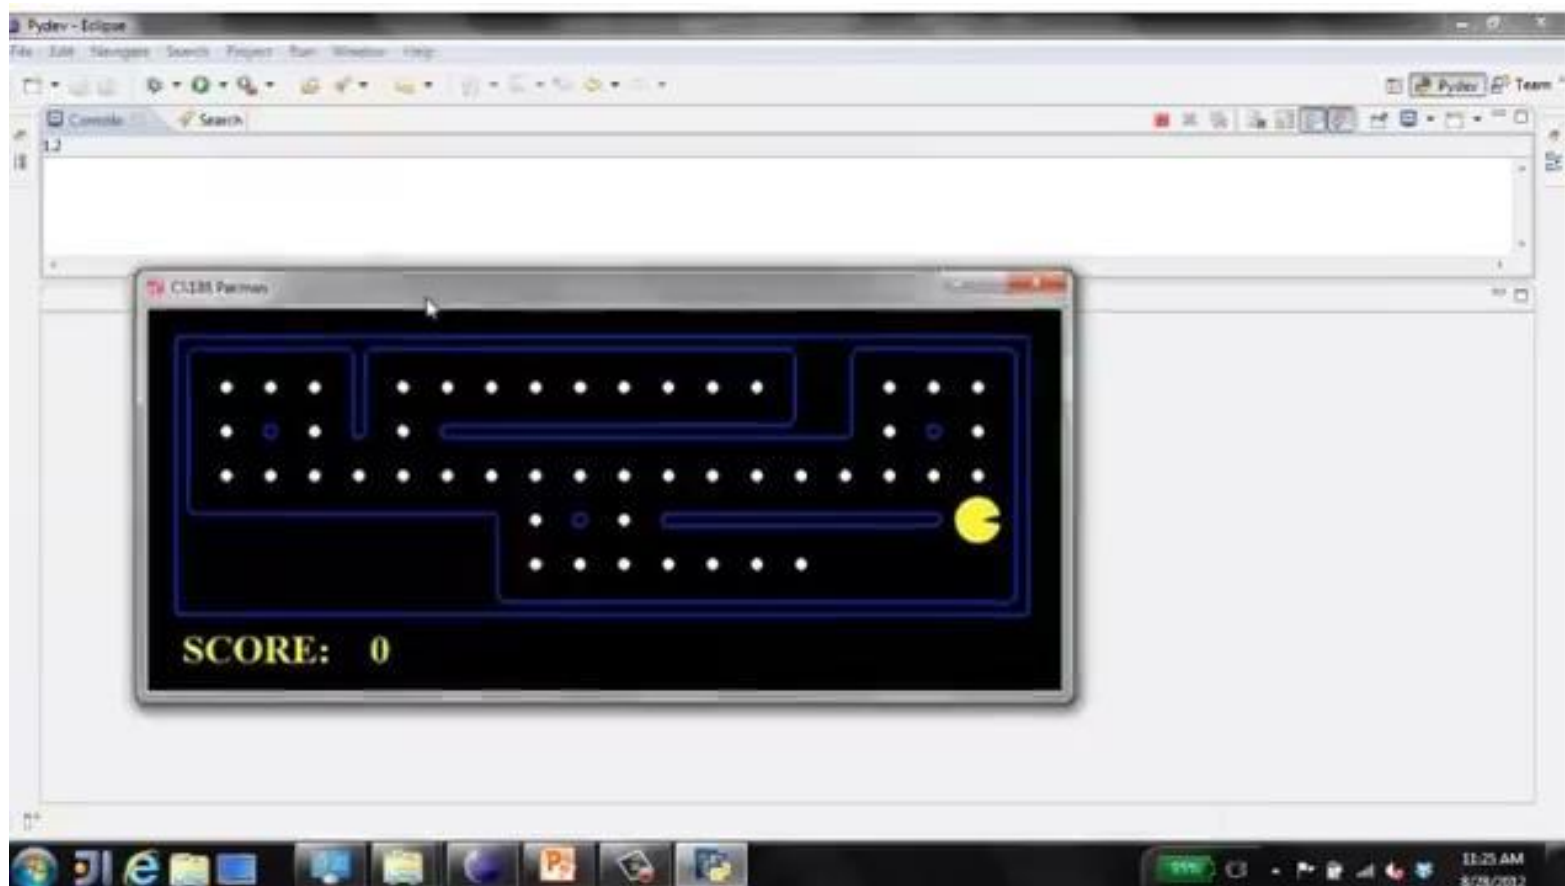


# Replanning Video

---




# Mastermind Video





# Our Goal: Help Pac-man finds its way!



Search: breadth- first, depth-first,  
uniform cost search.

Heuristic Search: Best-first,  $A^*$

# State-Space Search Problems

---

**General problem:** Find a path from a start state to a goal state given:

- A goal test: Tests if a given state is a goal state
- A successor function (transition model): Given a state and action, generate successor state

## Variants:

- Find any path vs. a least-cost path (if each step has a different cost i.e. a “step-cost”)
- Goal is completely specified, task is to find a path or least-cost path (i.e., Route planning)
- Path doesn't matter, only finding the goal state – 8 puzzle, N queens, Rubik's cube


# Search Problems

- A **search problem** consists of:

- A state space



- A successor function  
(with actions, costs)




- A start state and a goal test

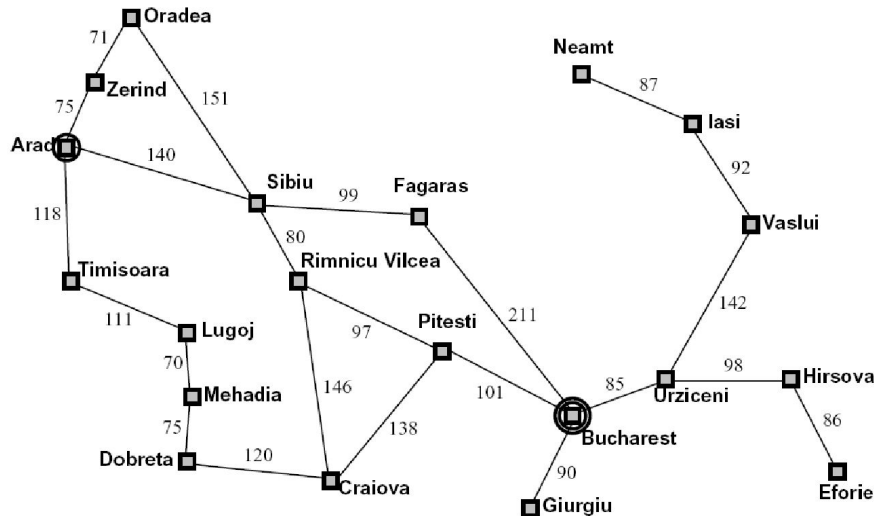
- A **solution** is a sequence of actions (a plan) which transforms the start state to a goal state

# Search Problems Are Models

---




# Example: Traveling in Romania



- State space:
  - Cities
- Successor function:
  - Roads: Go to adjacent city with cost = distance
- Start state:
  - Arad
- Goal test:
  - Is state == Bucharest?
- Solution?

# What's in a State Space?

The **world state** includes every last detail of the environment



A **search state** keeps only the details needed for planning (abstraction)

- **Problem: Pathing**


- States:  $(x,y)$  location
- Actions: NSEW
- Successor: update location only
- Goal test: is  $(x,y)=END$

- **Problem: Eat-All-Dots**


- States:  $\{(x,y), \text{dot booleans}\}$
- Actions: NSEW
- Successor: update location and possibly a dot boolean
- Goal test: dots all false

# State Space Sizes?

- World state:
  - Agent positions: 120
  - Food count: 30
  - Ghost positions: 12
  - Agent facing: NSEW
- How many
  - World states?  
 $120 \times (2^{30}) \times (12^2) \times 4$
  - States for pathing?  
120
  - States for eat-all-dots?  
 $120 \times (2^{30})$



# Quiz: Safe Passage




- Problem: eat all dots while keeping the ghosts perma-scared
- What does the state space have to specify?
  - (agent position, dot booleans, power pellet booleans, remaining scared time)




# State Space Graphs and Search Trees

---




# State Space Graphs

- State space graph: A mathematical representation of a search problem
  - Nodes are (abstracted) world configurations
  - Arcs represent successors (action results)
  - The goal test is a set of goal nodes (maybe only one)
- In a state space graph, each state occurs only once!
- We can rarely build this full graph in memory (it's too big), but it's a useful idea




# State Space Graphs

- In a search graph, each state occurs only once!
- We can rarely build this full graph in memory (it's too big), but it's a useful idea



*Tiny search graph for a tiny search problem*


# Search Trees



- A search tree:
  - A “what if” tree of plans and their outcomes
  - The start state is the root node
  - Children correspond to successors
  - Nodes show states, but correspond to PLANS that achieve those states
  - For most problems, we can never actually build the whole tree

# State Space Graphs vs. Search Trees


## State Space Graph



*Each NODE in in the search tree is an entire PATH in the state space graph.*

*We construct both on demand – and we construct as little as possible.*


## Search Tree



# Quiz: State Space Graphs vs. Search Trees

---

Consider this 4-state graph:




How big is its search tree (from S)?



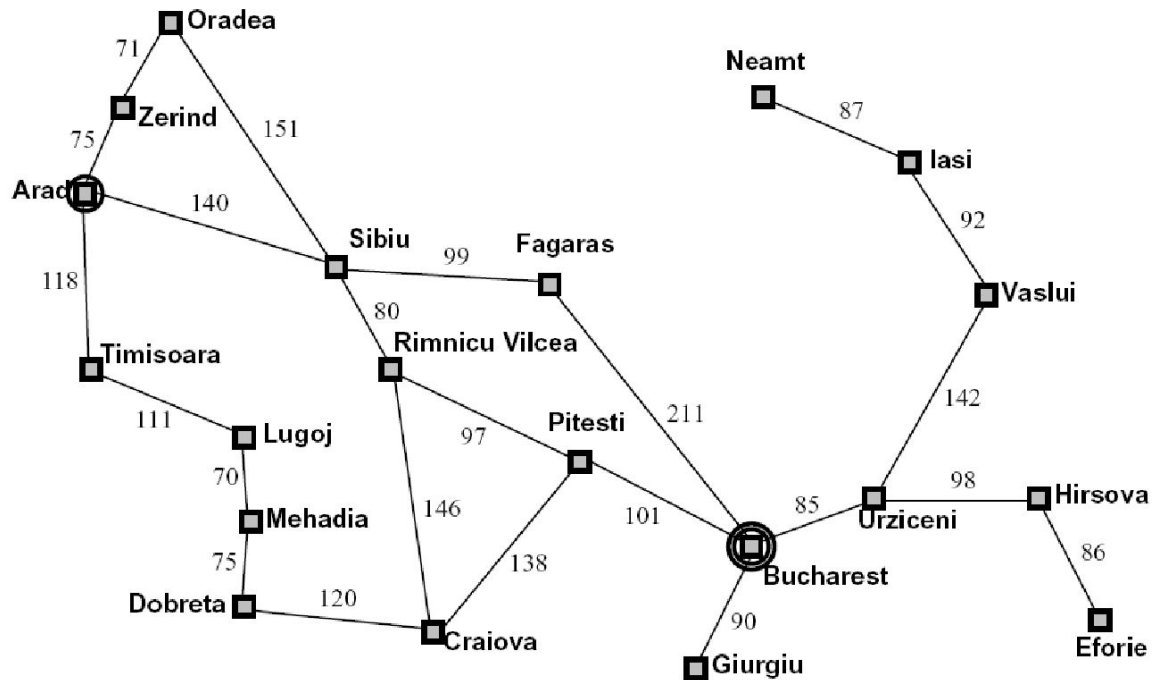
Important: Lots of repeated structure in the search tree!

# Tree Search

---




# Search Example: Romania





# Searching with a Search Tree




- Search:
  - Expand out potential plans (tree nodes)
  - Maintain a **fringe** of partial plans under consideration
  - Try to expand as few tree nodes as possible

# General Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```

- Important ideas:
  - Fringe
  - Expansion
  - Exploration strategy
- Main question: which fringe nodes to explore?

# Example: Tree Search



# Depth-First Search

---





# Depth-First Search

*Strategy: expand a  
deepest node first*


*Implementation:*

*Fringe is a LIFO stack*




# Search Algorithm Properties

---




# Search Algorithm Properties

- Complete: Guaranteed to find a solution if one exists?
- Optimal: Guaranteed to find the least cost path?
- Time complexity?
- Space complexity?
- Cartoon of search tree:
  - $b$  is the branching factor
  - $m$  is the maximum depth
  - solutions at various depths
- Number of nodes in entire tree?
  - $1 + b + b^2 + \dots + b^m = O(b^m)$



# Depth-First Search (DFS) Properties


- What nodes DFS expand?
  - Some left prefix of the tree.
  - Could process the whole tree!
  - If  $m$  is finite, takes time  $O(b^m)$
- How much space does the fringe take?
  - Only has siblings on path to root, so  $O(bm)$
- Is it complete?
  - $m$  could be infinite, so only if we prevent cycles (more later)
- Is it optimal?





# Breadth-First Search



---



# Breadth-First Search


*Strategy: expand a shallowest node first*

*Implementation:  
Fringe is a FIFO queue*




# Breadth-First Search (BFS) Properties

- What nodes does BFS expand?
  - Processes all nodes above shallowest solution
  - Let depth of shallowest solution be  $s$
  - Search takes time  $O(b^s)$
- How much space does the fringe take?
  - Has roughly the last tier, so  $O(b^s)$
- Is it complete?
  - $s$  must be finite if a solution exists, so yes!
- Is it optimal?
  - Only if costs are all 1 (more on costs later)



# Quiz: DFS vs BFS

---



# Quiz: DFS vs BFS


---

- When will BFS outperform DFS?
- When will DFS outperform BFS?


[Demo: dfs/bfs maze water (L2D6)]

# Iterative Deepening

- Idea: get DFS's space advantage with BFS's time / shallow-solution advantages
  - Run a DFS with depth limit 1. If no solution...
  - Run a DFS with depth limit 2. If no solution...
  - Run a DFS with depth limit 3. ....
- Isn't that wastefully redundant?
  - Generally most work happens in the lowest level searched, so not so bad!




# Cost-Sensitive Search



BFS finds the shortest path in terms of number of actions.  
It does not find the least-cost path. We will now cover  
a similar algorithm which does find the least-cost path.

# Uniform Cost Search

---







# Uniform Cost Search


Strategy: expand a  
cheapest node first:

Fringe is a priority queue  
(priority: cumulative cost)





# Uniform Cost Search (UCS) Properties

- What nodes does UCS expand?
  - Processes all nodes with cost less than cheapest solution!
  - If that solution costs  $C^*$  and arcs cost at least  $\varepsilon$ , then the “effective depth” is roughly  $C^*/\varepsilon$
  - Takes time  $O(b^{C^*/\varepsilon})$  (exponential in effective depth)
- How much space does the fringe take?
  - Has roughly the last tier, so  $O(b^{C^*/\varepsilon})$
- Is it complete?
  - Assuming best solution has a finite cost and minimum arc cost is positive, yes!
- Is it optimal?
  - Yes! (Proof next lecture via A\*)



# Uniform Cost Issues

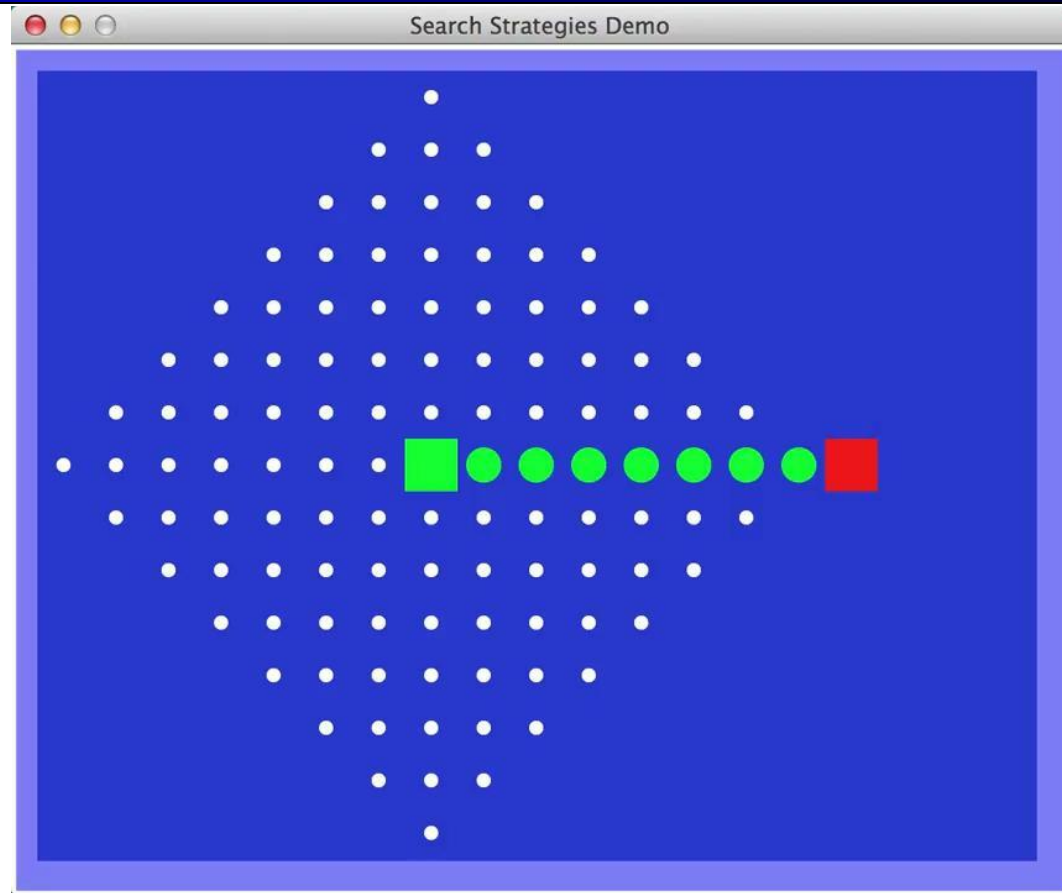
- Remember: UCS explores increasing cost contours
- The good: UCS is complete and optimal!
- The bad:
  - Explores options in every “direction”
  - No information about goal location
- We'll fix that soon!




[Demo: empty grid UCS (L2D5)]

[Demo: maze with deep/shallow water  
DFS/BFS/UCS (L2D7)]


# Video of Demo Empty UCS




## Video of Demo Maze with Deep/Shallow Water --- DFS, BFS, or UCS? (part 1)



## Video of Demo Maze with Deep/Shallow Water --- DFS, BFS, or UCS? (part 2)




## Video of Demo Maze with Deep/Shallow Water --- DFS, BFS, or UCS? (part 3)



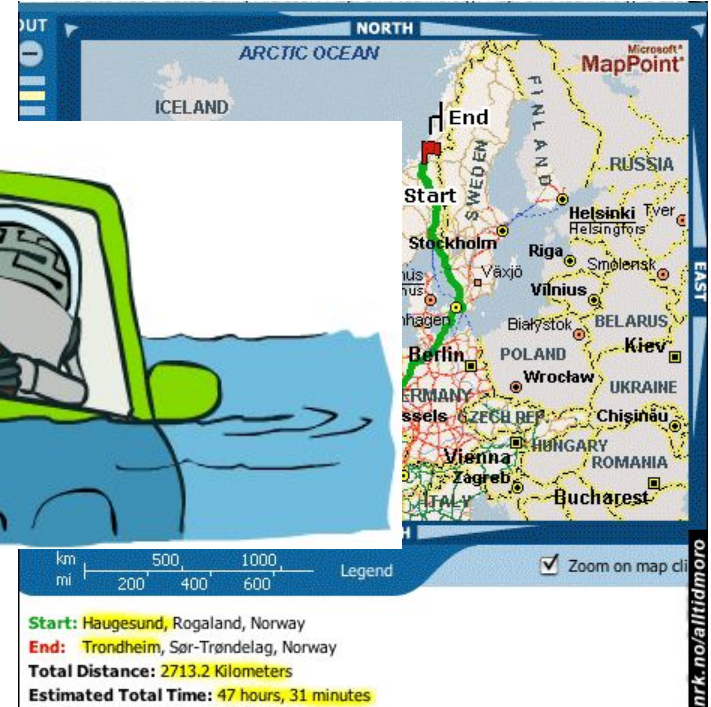
# Search and Models

- Search operates over models of the world
  - The agent doesn't actually try all the plans out in the real world!
  - Planning is all “in simulation”
  - Your search is only as good as your models...





# Search Gone Wrong?



# Some Hints for P1

---

- Graph search is almost always better than tree search (when not?)
- Implement your closed list as a dict or set!
- Nodes are conceptually paths, but better to represent with a state, cost, last action, and reference to the parent node

# Implementation

<https://inst.eecs.berkeley.edu/~cs188/fa24/projects/proj1/>

```
python pacman.py
```

```
python pacman.py -h
```

```
python pacman.py -l tinyMaze -p SearchAgent -a  
fn=tinyMazeSearch
```

## Testando SearchAgent:


```
python pacman.py -l tinyMaze -p SearchAgent -a  
fn=tinyMazeSearch
```

- Padrão é rodar DFS
- Encontrar a posição (1,1) → PositionSearchProblem

```
def __init__(self, fn='depthFirstSearch', prob='PositionSearchProblem', heuristic='nullHeuristic'):
```



# Main files

<code>search.py</code>	Where all of your search algorithms will reside.
<code>searchAgents.py</code>	Where all of your search-based agents will reside.
Files you might want to look at:	
<code>pacman.py</code>	The main file that runs Pacman games. This file describes a Pacman GameState type, which you use in this project.
<code>game.py</code>	The logic behind how the Pacman world works. This file describes several supporting types like AgentState, Agent, Direction, and Grid.
<code>util.py</code>	Useful data structures for implementing search algorithms.




Código: <https://inst.eecs.berkeley.edu/~cs188/fa24/assets/projects/search.zip>

# TinyMaze Layout



# Qual o tamanho da árvore?






?

```

def depthFirstSearch(problem):
    fringe = Stack()
    fringe.push(problem.getStartState())
    visited = []
    path=[]
    pathToCurrent=Stack()
    currState = fringe.pop()
    while not problem.isGoalState(currState):
        if currState not in visited:
            visited.append(currState)
            successors = problem.getSuccessors(currState)
            for child,direction,cost in successors:
                fringe.push(child)
                tempPath = path + [direction]
                pathToCurrent.push(tempPath)
            currState = fringe.pop()
            path = pathToCurrent.pop()
    return path

```



# The 8-puzzle

---

**Start**

1	2	3
8		4
7	6	5



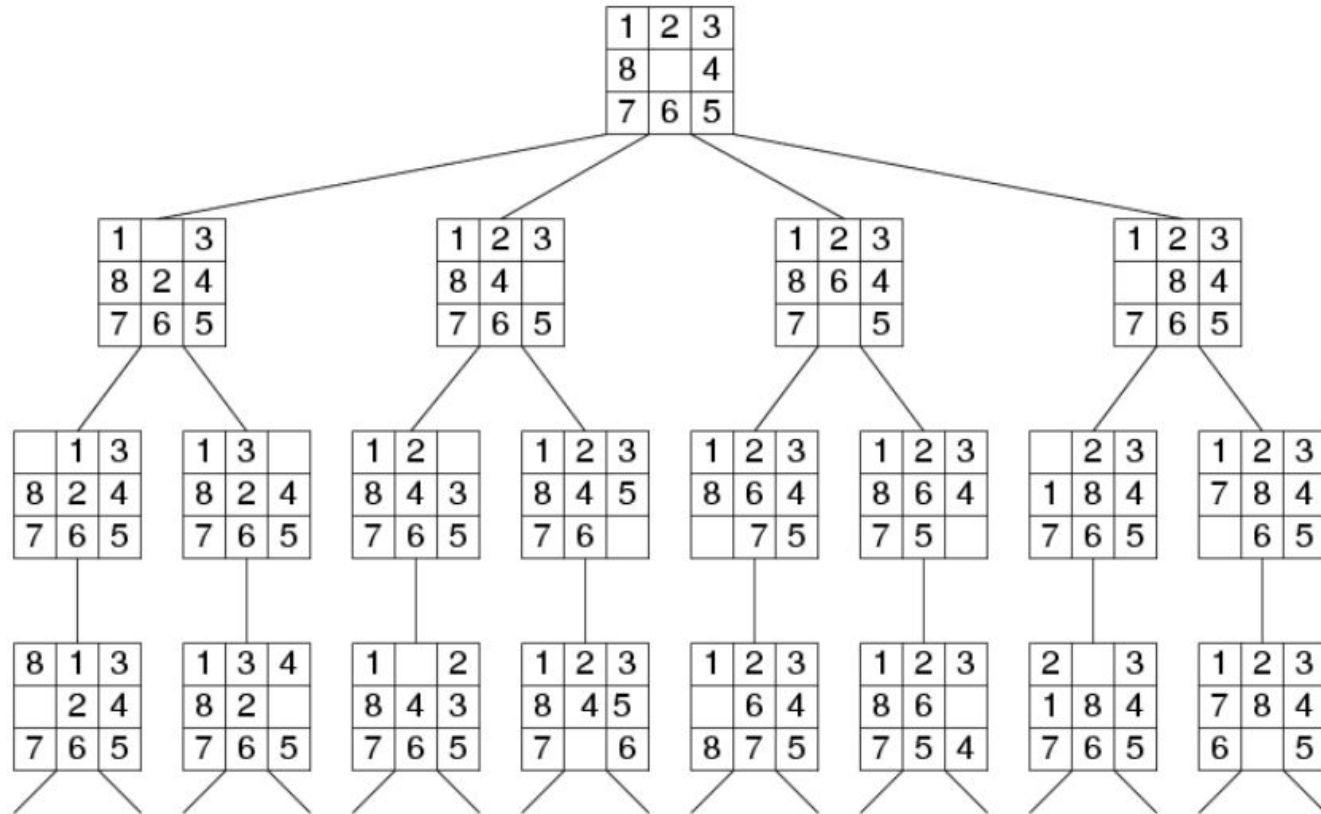
**Goal**

1	2	3
4	5	6
7	8	

Como modelar isto?



# Tree Structure



# Facebook friends



Find the shortest chain of Facebook friends that goes from Person A to Person B

# Robotics



Commercial



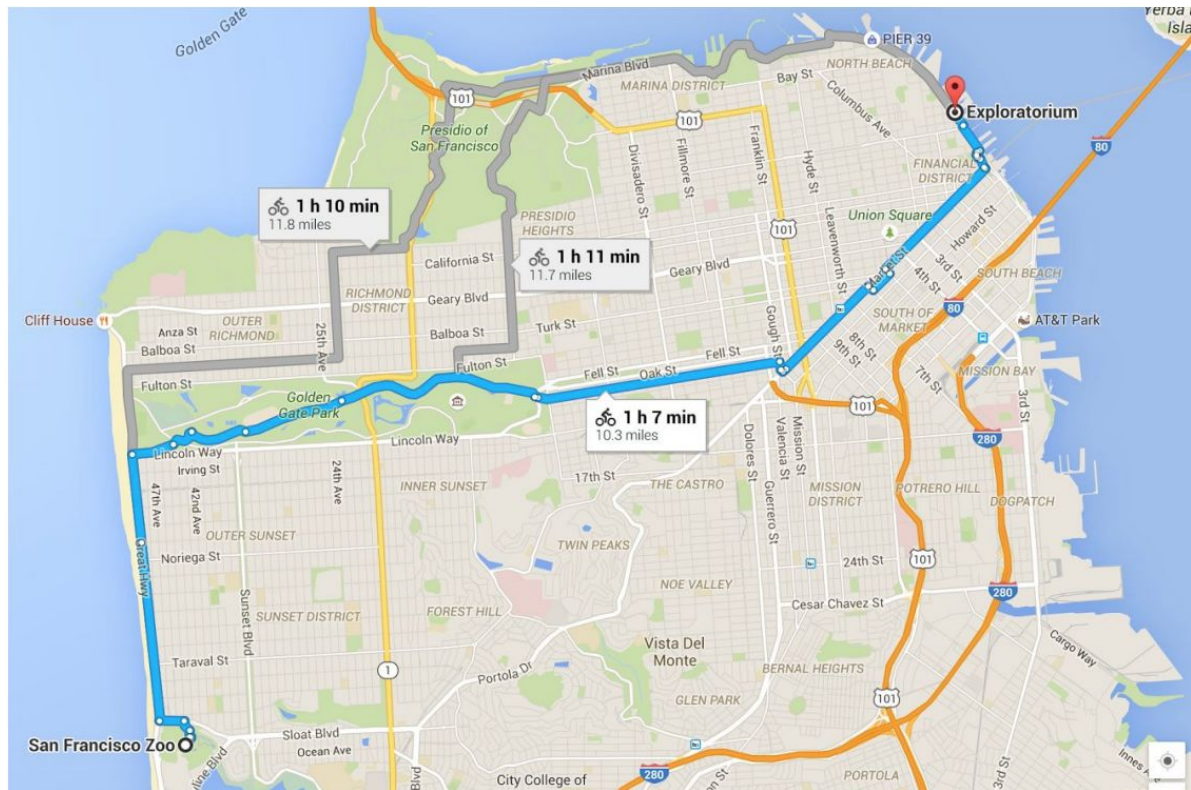
Search & Rescue



Domestic

Source: <https://cs.stanford.edu/people/abisee/gs.pdf>

# Route Planning






# Time and Space Complexity

Time and memory usage when  $b = 10$ :

<b>solution depth</b>	<b>nodes considered</b>	<b>time</b>	<b>memory</b>
0	1	1 millisecond	100 bytes
4	11,111	11 seconds	1 megabyte
8	$10^8$	31 hours	11 gigabytes
10	$10^{10}$	128 days	<b>1 terabyte</b>
12	$10^{12}$	35 years	111 terabytes

# Exercise 1



Node	$h_1$	$h_2$
A	9.5	10
B	9	12
C	8	10
D	7	8
E	1.5	1
F	4	4.5
G	0	0

Consider the state space graph shown above. A is the start state and G is the goal state. The costs for each edge are shown on the graph. Each edge can be traversed in both directions. Note that the heuristic  $h_1$  is consistent but the heuristic  $h_2$  is not consistent.

**(a) Possible paths returned**

For each of the following graph search strategies (*do not answer for tree search*), mark which, if any, of the listed paths it could return. Note that for some search strategies the specific path returned might depend on tie-breaking behavior. In any such cases, make sure to mark *all* paths that could be returned under some tie-breaking scheme.

Search Algorithm	A-B-D-G	A-C-D-G	A-B-C-D-F-G
Depth first search			
Breadth first search			
Uniform cost search			
A* search with heuristic $h_1$			
A* search with heuristic $h_2$			



# Exercise 2

---

Try running our 8-puzzle solver on the initial state shown at right!

	8	7
6	5	4
3	2	1



# Hanoi Tower

---



Implementation using Depth First Search

# Homework - 25%

---

1. Python DFS, BFS, UCS (Berkeley framework)
2. Exercise 1 - by hand
3. Hanoi Tower - Python

**Deadline:** May, 27rd, 2025